Edit Distance

1 The Edit Distance Problem

Definition of a ref-word

Given a finite set of variables $V \subseteq \text{SVars}$ we define the alphabet of ref words as: $\Gamma_V := \{x\vdash, \dashv x\}$. And given an alphabet Σ such that $\Sigma \cap \Gamma_V = \varnothing$ we can define the set of ref words over Σ and V as: $\mathbf{r} \in (\Sigma \cup \Gamma_V)^*$. Next, a ref-word is valid if and only if, every occurance of a variable in the ref-word is opened exactly once and closed afterwards, exactly once.

Functions on ref-words

We can define the projection of a ref word over a set S, $r \upharpoonright S$, recursively as:[1]

1. $r \in S \rightarrow r \upharpoonright S = r$
2. $r \notin S \rightarrow r \upharpoonright S = \epsilon$
3. $(r_1 \cdot r_2) \upharpoonright S = (r_1 \upharpoonright S) \cdot (r_2 \upharpoonright S)$

Vars(r) is the set of variables $x \in V$ that occurs in the ref-word:

$$\text{Vars}(r) := \left\{x \in V \;\middle|\; \exists r_x^{pre}, r_x, r_x^{post} \in (\Sigma \cup \Gamma_V)^* \text{ such that } r = r_x^{pre} \cdot x\vdash \cdot r_x \cdot \dashv x \cdot r_x^{post})\right\} \tag{1}$$

$\text{tup}(r)$ are the positions each ref-word is referencing, and is defined as:

$$\text{tup}(r) := \{x \mapsto [i_x, j_x\rangle \mid x \in Vars(R), i_x = |r_x^{pre} \upharpoonright \Sigma|, j_x = i_x + |r_x \upharpoonright \Sigma|\} \tag{2}$$

Postulate:

$$\text{valid}(r) \rightarrow |\text{tup}(r)| = |\text{Vars}(r)| \tag{3}$$

Definition of ref-word tuple

Distance between two ref words

Next, given two ref words $r_1, r_2 \in (\Sigma \cup \Gamma_V)$ and a distance function $\mathbf{d} : \Sigma^* \times \Sigma^* \rightarrow \mathbb{R}$ the distance $\mathbf{d_\Sigma}$ between r_1 and r_2 is defined as: $\mathbf{d_\Sigma}(r_1, r_2) = \mathbf{d}(r_1 \upharpoonright \Sigma, r_2 \upharpoonright \Sigma)$ and the distance $\mathbf{d_\Gamma}$ is defined extending the distance function $\mathbf{b}$ over $\Sigma \cup \Gamma_V$, treating letters in Γ_V equivalently as words in Σ.

Ref-word distance languages

Given a ref-word language $R \subseteq (\Sigma \cup \Gamma_V)^*$ and a distance $k \in \mathbb{R}$, the k-distance ref-word language is defined as:

$$R \pm k = \{r \in (\Sigma \cup \Gamma_V)^* \mid \text{valid(r)}, \exists r'(r' \in R \wedge \mathbf{d_\Sigma}(r, r') = \mathbf{d_\Gamma}(r, r') \leq k\} \tag{4}$$

Given a document d, the spanner over a ref-word language R is:

$$[\![R]\!]_d = \{\text{tup}(r) \mid r \upharpoonright \Sigma = d, \exists r \in R\} \tag{5}$$

[1] In the paper (Doleschal, 2021) this operation is defined for Σ as $\text{doc}(\sigma)$

Theorem 1. *If the distance function* $\mathbf{d}$ *is a metric, then:*

$$(R \pm n) \pm k = R \pm (n + k) \tag{6}$$

Proof. Given $n, k \in \mathbb{R}$, we want to prove that: $(R \pm n) \pm k = R \pm (n + k)$. By definition, we have that:

$$(R \pm n) \pm k = \{r \in (\Sigma \cup \Gamma_V)^* \mid \text{valid(r)}, \exists r'(r' \in R \pm n \wedge \mathbf{d_\Sigma}(r, r') \leq k)\}$$

First we prove that $(R \pm n) \pm k \subseteq R \pm (n + k)$. By contradiction let's assume there exists an element $r_1 \in (R \pm n) \pm k$ such that $r_1 \notin R \pm (n + k)$. By the previous definition, we have that $\exists r' \in R \pm n$ such that $\mathbf{d_\Sigma}(r_1, r') \leq k$. By definition of $R \pm (n)$ we have that for any $r \in R$, $\mathbf{d_\Sigma}(r', r) \leq n$. Next, by the definition of $R \pm (n + k)$, and our suposition we have that $\mathbf{d_\Sigma}(r_1, r) > n + k$ which contradicts the triangle inequality. The proof that $R \pm (n + k) \subseteq (R \pm n) \pm k$ uses this same argument.

□

Variable-set automaton over ref words (VSet-automaton)[2]

Definition 1. *A VSet-automaton is a sextuple* $A := (\Sigma, V, Q, q_0, Q_F, \delta)$

- Σ*: Alphabet symbols*
- V*: Finite set of variables*
- Q*: Finite set of states*
- $q_0 \in Q$*: Initial state*
- $Q_F \subseteq Q$*: Set of final states*
- $\delta : Q \times (\Sigma \cup \{\epsilon\} \cup \Gamma_V) \to 2^Q$*: Transition function*
  - $\Gamma_V := \{x \vdash, \dashv x \mid x \in V\}$
  - 2^Q*: power set of* Q

Ref-word language

The ref-word language of A is: $\mathcal{R}(A) = \mathcal{R}^0(A) = \{r \in \mathcal{L}(A) \subseteq (\Sigma \cup \Gamma_V)^* \mid r$ is accepted by the $\epsilon -$ NFA A$\}$. This is direct from interpreting A as an $\epsilon - NFA$.

Run of a VSet-automaton over a ref-word

Given a ref-word $r = \sigma_1 \cdots \sigma_n$, the run ρ of A is the sequence:

$$\rho := q_0 \xrightarrow{\sigma_1} q_1 \cdots q_{n-1} \xrightarrow{\sigma_n} q_n \tag{7}$$

Where $\forall i \in [0, n)\,(q_{i+1} \in \delta(q_i, \sigma_{i+1}))$ and $q_n \in Q_F$
From previous publications we know that $r \in \mathcal{R}(A)$ if and only if there is a run ρ of A on r.

[2] (Doleschal, 2021)

Distance automaton

Given a VSet-automaton A we can define, under Levenshtein distance, the automaton $A\pm1 := (\Sigma, V, Q', q'_0, Q'_F, \delta')$ Where

- $Q' = \{q_1, ..., q_{|Q|}\} \cup \left\{q^1_1, ..., q^1_{|Q|}\right\}$ Where there exists two biyective functions:

  1. $f : Q \to \{q_1, ..., q_{|Q|}\}$
  2. $f' : Q \to \left\{q^1_1, ..., q^1_{|Q|}\right\}$

  And two biyective functions F and F' that map f and f' respectively to sets.

- $q'_0 = f(q_0)$
- $Q'_F = Q_F \cup \{q^1_i \mid q_j \in Q_F \wedge f'(q_j) = q^1_i\}$
- The function δ' is defined by:

$$\delta'(q_i, e) = \begin{cases} F(\delta(f^{-1}(q_i), e)) & e \in \Gamma_V \\ F(\delta(f^{-1}(q_i), e)) \cup q^1_i \cup \bigcup\limits_{a\in\Sigma} F'(\delta(f^{-1}(q_i), a)) & e.o.c \end{cases}$$

$$\delta'(q^1_i, e) = F'(\delta(f'^{-1}(q^1_i), e))$$

Definition 2. *A ref-word language R is sequential if every ref-word $r \in R$ is valid.*

Lemma 1. *All runs of $\mathcal{R}(A \pm 1)$ have one of the following structures:*

1.
$$\rho_{A\pm1} = f(q_0) \xrightarrow{c_1} f(\phi_1) \cdot\cdot \xrightarrow{c_{a-1}} f(\phi_{a-1}) \xrightarrow{c_a} f'(\phi_{a-1}) \cdot\cdot \xrightarrow{c_n} f'(\phi_{n-1})$$

2.
$$\rho_{A\pm1} = f(q_0) \xrightarrow{c_1} f(\phi_1) \cdot\cdot \xrightarrow{c_{a-1}} f(\phi_a) \xrightarrow{c_a} f'(\phi_{a+1}) \xrightarrow{c_{a+1}} f'(\phi_{a+2}) \cdot\cdot \xrightarrow{c_n} f'(\phi_n)$$

3.
$$\rho_{A\pm1} = f(q_0) \xrightarrow{c_1} f(\phi_1) \cdot\cdot \xrightarrow{c_{a-1}} f(\phi_a) \xrightarrow{\epsilon} f'(\phi_{a+1}) \xrightarrow{c_a} f'(\phi_{a+2}) \cdot\cdot \xrightarrow{c_n} f'(\phi_{n+1})$$

4.
$$\rho_{A\pm1} = f(q_0) \xrightarrow{c_1} f(\phi_1) \xrightarrow{c_2} \cdots \xrightarrow{c_n} f(\phi_n)$$

Proof. This is direct from the form of the transitions under the definition of $A \pm 1$, and noticing that there is no transition from $f'(q)$ to $f(q)$. □

Theorem 2. *if $\mathcal{R}(A)$ if sequential, then $\mathcal{R}(A \pm 1)$. is sequential.*

Proof. For purposes of contradiction let's assume that there exists a ref-word $r = c_1 \cdots c_n \in A \pm 1$ that is not valid. Then there must exist a run $\rho_{A\pm1}$ on r, and furthermore, by lemma 1, $\rho_{A\pm1}$ it can have only 4 structures. We will prove by enumeration that all these structures lead to a contradicton.

1. **Insertion**: By definition of the transitions of $A \pm 1$, we have that $c_1 \cdots c_{a-1} \cdot c_{a+1} \cdots c_n \in \mathcal{R}(A)$. This ref-word is valid by the premise of the theorem. Next, c_a cannot be in Γ_V, and therefore we have a contradiction, because variables in all other characters open exactly once, and close exactly once only after they are opened by the definition of a valid ref-word.
2. **Substitution, Elimination** Follow the same argument, in that the word changed is not in Γ_V.
3. **No modification** Trivial.

□

Theorem 3. *Given a VSet-automaton with a sequential ref-word language $\mathcal{R}(A)$, using Levenshtein distance we obtain that $\mathcal{R}(A \pm 1) = \mathcal{R}(A) \pm 1$.*

Proof. This is equivalent to proving that, given a ref-word r, $r \in \mathcal{R}(A \pm 1) \leftrightarrow r \in \mathcal{R}(A) \pm 1$. Therefore this is a two part proof.
First let's assume that $r \in \mathcal{R}(A \pm 1)$. In that case we know that there must exist a run ρ of $A \pm 1$ on r. r is valid due to the theorem 2. By Lemma 1 we have four cases for runs of $\mathcal{R}(a)$, We will now prove by enumeration on these cases:

1. **Insertion:** From the definition of $A \pm 1$, there must be a transition from ϕ_a to ϕ_{a+1} using the letter $a+1$. Therefore, there is a run:

   $$q_0 \xrightarrow{c_1} \cdots \xrightarrow{c_{a-1}} \phi_{a-1} \xrightarrow{c_{a+1}} \phi_a \cdot\cdot \xrightarrow{c_n} \phi_{n-1}$$

   And therefore, $c_1 \cdots c_a \cdot c_{a+2} \cdots c_n \in \mathcal{R}(A)$. And because the extended Levenshtein distance is of 1, and r is valid, then $r \in \mathcal{R}(A) \pm 1$

2. **Substitution, Elimination**: Follow the same argument, changing the sequence into a word and therefore, using the definitions of transitions arriving to a word in $\mathcal{R}(A)$.

3. **No modification**: In this case, we can obtain a run over the same word in $\mathcal{R}(A)$, and therefore, by definition of $\mathcal{R}(A) \pm 1$ we obtain that $r \in \mathcal{R}(A)$

Next, let's assume that $r \in \mathcal{R}(A) \pm 1$. In that case, r is valid, and there exists $r' \in R$ such that using Levenshtein distance: $\mathbf{d_\Sigma}(r, r') = \mathbf{d_\Gamma}(r, r') \leq 1$. Since Levenshtein distance is discrete, there are two cases:

1. $\mathbf{d_\Sigma}(r, r') = 0$. Then it is clear that $r \in \mathcal{R}(A \pm 1)$ since a subset of $A \pm 1$ forms an isomorphism with A.

2. $\mathbf{d_\Sigma}(r, r') = 1$. In this case, because of the structure of Levenshtein's distance, there are three possible cases:

   - **Insertion**. In this case the word r looks like: $c_1 \cdots c_i \cdot c_{inserted} \cdot c_{i+1} \cdots c_n$. Furthermore, $c_{inserted} \in \Sigma$ because if not, $\mathbf{d_\Gamma}(r, r') > \mathbf{d_\Sigma}(r, r')$. Next, since $c_1 \cdots c_n \in \mathcal{R}(A)$ this word is accepted by the automaton $A \pm 1$ by the definitions of the transitions.
   - **Substitution / Elimination** Same arguments as Insertion.

□

Theorem 4. *For any sequential automatons A there exists an automaton B such that $\mathcal{R}(B) = \mathcal{R}(\mathcal{A}) \pm k$ for all $k \in \mathbb{N}$ under Levenshtein distance.*

Proof. Trivial proof using theorems 1 and 3, and induction over k. □

2 Alternative Definitions

2.1 Navarro 2001

The approximate search problem can therefore be stated as follows. Given a text $\mathcal{T}$, a regular expression $\mathcal{E}$, an edit distance between strings $d()$, and a distance threshold k, find all the text positions that start an approximate occurrence of E in T , that is, compute the set:

$$\{i, \exists j, \exists S \in L(E), d(T_{i \ldots j}, S) \leq k\}$$

2.2 Brainstorming - Vrgoc

$$\{i, j, k' \mid \exists S \in L(E), d(T_{i \ldots j}, S) = k' \leq k\}$$